Malloc Bootcamp

Kashish, Adithi, Jack

March 19, 2020

Agenda

- Conceptual Overview
    - Explicit List
    - Segregated list
    - Splitting, coalescing
    - Hints on hints
- Advanced debugging with GDB
    - Fun GDB tricks
- Writing a good heap checker
- Appendix

Conceptual Outline

Dynamic Memory Allocation

- Used when
  - we don't know at compile-time how much memory we will need
  - when a particular chunk of memory is not needed for the entire run
    - lets us reuse that memory for storing other things
- Important terms:
  - malloc/calloc/realloc/free
  - mem_sbrk
  - payload
  - fragmentation (external vs internal)
  - Splitting / coalescing

Anonymous Structs/Unions

Same idea with unions. For the difference between unions and structs, refer to the C bootcamp slides.

struct name

```
struct A {
    int x;
    struct B {
        int y;
        float z;
    } my_b;
} my_a;
```

member name

```
struct A {
    int x;
    struct B {
        int y;
        float z;
    };
} my_a;
```

- What is the type of `x`?   **int**
- How do we access `x` or `my_a`?   **my_a.x**
- What is the type of `my_b`?   **struct B**   **struct B**
- How do we access `y` of `my_a`?   **my_a.my_b.y**   **my_a.y**

5

```c
/** @brief Represents the header and payload of one block in the heap */
typedef struct block {
    /** @brief Header contains size + allocation flag */
    word_t header;

    /**
     * @brief A pointer to the block payload.
     *
     * TODO: feel free to delete this comment once you've read it carefully.
     * We don't know what the size of the payload will be, so we will declare
     * it as a zero-length array, which is a GCC compiler extension. This will
     * allow us to obtain a pointer to the start of the payload.
     *
     * WARNING: A zero-length array must be the last element in a struct, so
     * there should not be any struct fields after it. For this lab, we will
     * allow you to include a zero-length array in a union, as long as the
     * union is the last field in its containing struct. However, this is
     * compiler-specific behavior and should be avoided in general.
     *
     * WARNING: DO NOT cast this pointer to/from other types! Instead, you
     * should use a union to alias this zero-length array with another struct,
     * in order to store additional types of data in the payload memory.
     */
    char payload[0];

    /*
     * TODO: delete or replace this comment once you've thought about it.
     * Why can't we declare the block footer here as part of the struct?
     * Why do we even have footers -- will the code work fine without them?
     * which functions actually use the data contained in footers?
     */
} block_t;
```
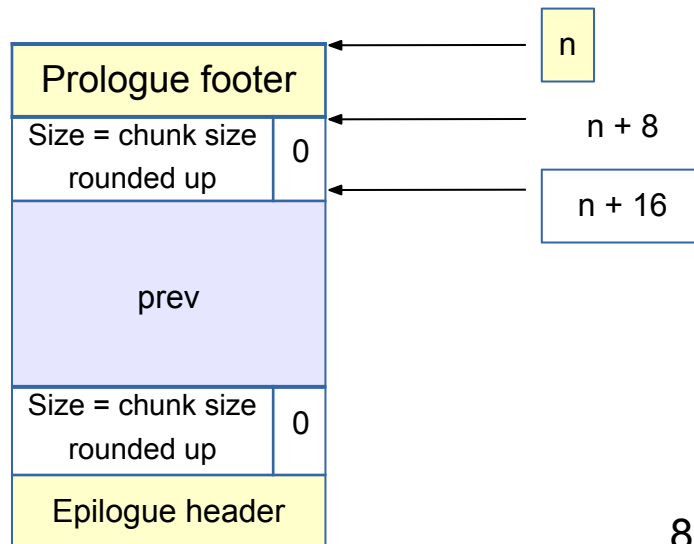
Zero-Length Arrays

```c
struct line {
    int length;
    char contents[0];
};

int main() {
    struct line my_line;
    printf("sizeof(contents) = %zu\n",    sizeof(L.contents));  // 0
    printf("sizeof(struct line) = %zu\n", sizeof(struct line)); // 4
}
```

- It's a GCC extension - not part of the C specification!
- Must be at the end of a struct / union
- It simply allows us to represent variable-length structures
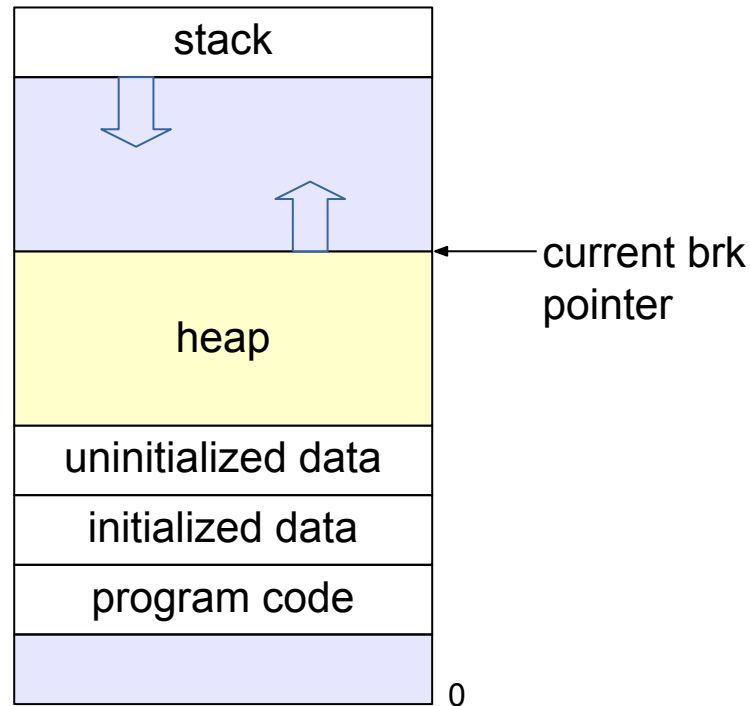- `sizeof` on a zero-length array returns zero

7

mm_init

- Why prologue footer and epilogue header?

- Payload must be 16-byte aligned

- But, the size of payload doesn't have to be a multiple of 16 - just the block does!

- Things malloc'd must be within the prologue and epilogue

n

Prologue footer	←
Epilogue header	←

n + 8

n

Prologue footer
Size = chunk size rounded up	0
prev	
Size = chunk size rounded up	0
Epilogue header	

n + 8

n + 16

8

If We Can't Find a Usable Free Block

- Assume an implicit list implementation
- Need to extend the heap
  - mem_sbrk()
    - sbrk(num_bytes) allocates space and returns pointer to start
    - sbrk(0) returns a pointer to the end of the current heap
- For speed, extend the heap by a little more than you need immediately
  - use what you need out of the new space, add the rest as a free block
 - What are some tradeoffs you can make?

stack
(arrows)
heap
uninitialized data
initialized data
program code

← current brk pointer

0

6

Tracking Blocks: Explicit List

- Maintain a list of *free* blocks instead of *all* blocks
    - means we need to store forward/backward pointers, not just sizes
    - we only track free blocks, so we can store the pointers in the payload area!
   - need to store size at end of block too, for coalescing

allocated block

size	1
payload and padding	
size	1

free block

size	0
next	
prev	
unused	
size	0

10
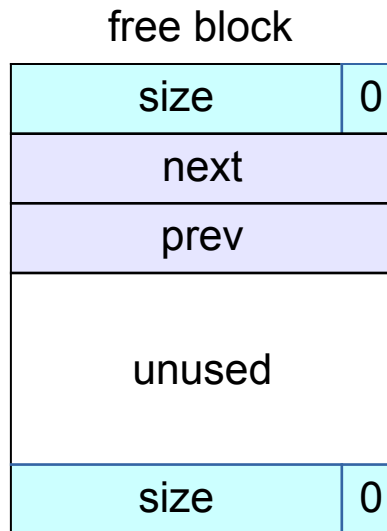
Splitting a Block

- If the block we find is larger than we need, split it and leave the remainder for a future allocation
  - explicit lists: correct previous and next pointers
  - Segregated lists: same as explicit
- When would we **not** split a block?

n	0
next	
prev	
n	0

→

m	1
payload	
m	1
n-m	0
next	
prev	
n-m	0

Coalescing Memory

- Combine adjacent blocks if both are free
  - **explicit lists**: look forward and backward **in the heap**, using block sizes, **not** next/prev

- Four cases:

block to
be freed

Allocated	Allocated	Free	Free
Allocated	Free	Allocated	Free

Coalescing Memory
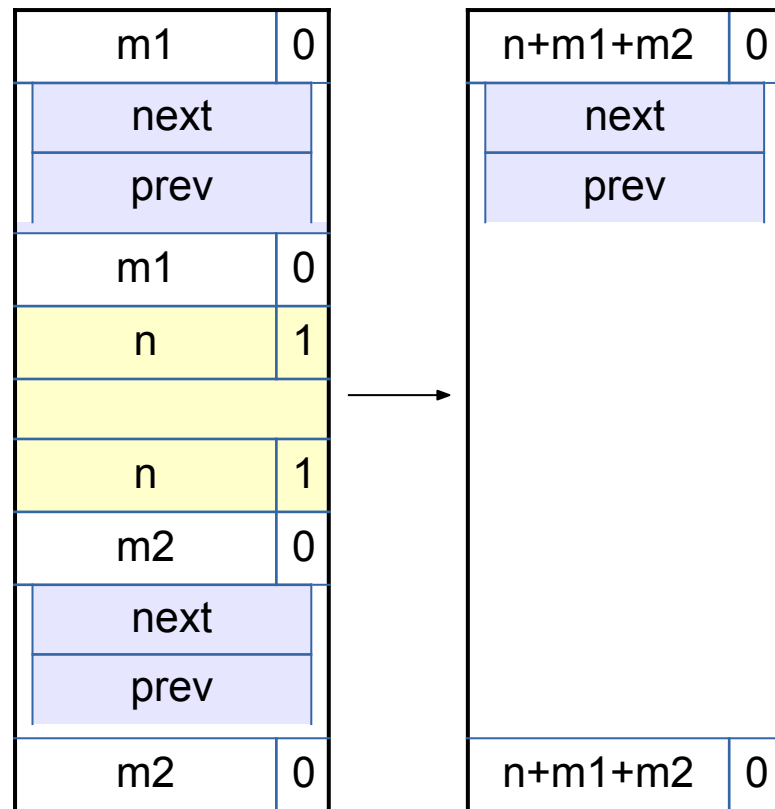
Design Considerations

- Finding a matching free block
  - First fit vs. next fit vs. best fit vs. "good enough" fit
  - continue searching for a closer fit after finding a big-enough free block?
- Free block ordering
  - LIFO, FIFO, or address-ordered?
- When to coalesce
  - while freeing a block or while searching for free memory?
- How much memory to request with sbrk()
  - larger requests save time in system calls but increase maximum memory use

14

Hints on hints

For the final, you must greatly increase the utilization and keep a high throughput.

- Reducing external fragmentation requires achieving something closer to best-fit allocated
    - Using a better fit algorithm
    - Combine with a better data structure that lets you run more complex algorithms

- Reducing internal fragmentation requires reducing data structure overhead and using a 'good' free block

15

Segregated Lists

- Multiple explicit lists where the free blocks are of a certain size range

- Increases throughput and raises probability of choosing a better-sized block

- Need to decide what size classes (only 128 bytes of stack space)
  - Diminishing returns
  - What do you do if you can't find something in the current size class?

- RootSizeClass1 -> free-block 1 -> free-block 2 -> free-block 3 ->

- RootSizeClass2 -> free-block 1 -> free-block 2 -> free-block 3 -> ...

- RootSizeClass3 -> free-block 1 -> free-block 2 -> free-block 3 -> ...

- ...

Modularity and Design

- Now you need to have more than one list
  - List operations are the same for all lists
    - Insert
    - Remove
  - Deciding which size class a block should go into
    - 14 if statements :(
    - A small **const** array of sizes + a loop :)
- It would be quite painful to maintain copy-pasted code
  - Abstractions are nice - it's what CS is all about!

Modularity and Design

- Make sure you have modular, extensible code
  - It will save you a lot of time spent debugging and style points
  - It will make you happy when you come back to your code
    - In 6 days when you start the final submission
    - Or in 6 hours if you're missing sleep - please get some rest!
  - It will make it easier to explain to students when you become a TA later :)
- Labs in this course are NOT meant to be done in one sitting
- Labs in this course are NOT meant to be done in 2-3 nights
- Plan ahead, leave plenty of time for **design**
  - Measure twice, cut once
- Take a break between sittings
  - Your brain can keep working subconsciously
  - Leave time for "aha!" moments

Coalescing Memory

- Combine adjacent blocks if both are free
  - **segregated lists**: look forward and back using block sizes, then
    - Use the **size** of the **coalesced** block to determine the proper list
      - What else might you need to do to maintain your seglists?
    - Insert into list using the insertion policy (LIFO, address-ordered, etc.)

- Four cases:

block to
be freed

Allocated	Allocated	Free	Free
Allocated	Free	Allocated	Free

19

Eliminate footers in allocated blocks

- Reduces internal fragmentation (increase utilization)
- Why do we need footers?
  - Coalescing blocks
  - What kind of blocks do we coalesce?
- Do we need to know the size of a block if we're not going to coalesce it?
- Based on that idea, can you design a method that helps you determine when to coalesce?
  - Hint: where could you store a little **bit** of extra information for each block?

m1	0
next	
prev	

m1	0
m2	1
payload	
m3	1
payload	

free blocks still have footers

allocated blocks don't have footers!

20

Coalescing Memory

- Combine adjacent blocks if both are free
  - **footerless:** if free, obtain info from footer then use next/prev

- Four cases:

block to
be freed

Allocated	Allocated	Free	Free
Allocated	Free	Allocated	Free

21

Decrease the minimum block size

- Reduces internal fragmentation (increase utilization)
- Currently, min block size is 32.
  - 8 byte header
  - 16 byte payload (or 2 8 byte pointers for free)
  - 8 byte footer
- If you just need to malloc(5), and the payload size is 16 bytes, you waste 11 bytes.
- Must manage free blocks that are too small to hold the pointers for a **doubly** linked free list

9 bytes are wasted!
How can we prevent this?

22

Debugging: GDB & The Almighty Heap Checker

What's better than printf? Using GDB

- Use GDB to determine where segfaults happen!
- **gdb mdriver** will open the malloc driver in gdb
  - Type run and your program will run until it hits the segfault!
- **step/next** - (abbrev. **s/n**) step to the next line of code
  - **next** steps over function calls
- **finish** - continue execution until end of current function, then break
- **print <expr>** - (abbrev. **p**) Prints **any C-like expression** (including results of function calls!)
  - Consider writing a heap printing function to use in GDB!
- **x <expr>** - Evaluate <expr> to obtain address, then examine memory at that address
  - **x /a <expr>** - formats as address
  - See **help p** and **help x** for information about more formats

24

Using GDB - Fun with frames

- **backtrace** - (abbrev. **bt**) print call stack up until current function
    - **backtrace full** - (abbrev. **bt full**) print local variables in each frame

(gdb) backtrace
#0 find_fit (...)
#1 mm_malloc (...)
#2 0x0000000000403352 in eval_mm_valid
(...) #3 run_tests (...)
#4 0x0000000000403c39 in main (...)

- **frame 1** - (abbrev. **f 1**) switch to mm_malloc's stack frame
    - Good for inspecting local variables of calling functions

Using GDB - Setting breakpoints/watchpoints

- **break mm_checkheap** - (abbrev. **b**) break on "mm_checkheap()"
  - **b mm.c:25** - break on line 25 of file "mm.c" - **very useful!**
- **b find_fit if size == 24** - break on function "find_fit()" if the local variable "size" is equal to 24 - "**conditional breakpoint**"

- **watch heap_listp** - (abbrev. **w**) break if value of "heap_listp" changes - "**watchpoint**"
- **w block == 0x80000010** - break if "block" is equal to this value
- **w *0x15213** - watch for changes at memory location 0x15213
  - Can be *very* slow

- **rwatch <thing>** - stop on reading a memory location
- **awatch <thing>** - stop on *any* memory access

What's better than GDB? Using CGDB!

- CGDB is just like GDB
  - but with COLOR and **SOURCE_CODE**
- Breaking at mm_malloc in GDB vs CGDB

Colors!

C Code!

GDB
terminal!

Using CGDB

- ## Initializes the same as GDB
  - Just write cgdb mdriver-dbg instead of gdb mdriver-dbg
- ## Source and gdb windows
  - To go from gdb (default) to source press esc (just like normal mode in vim!!!!!)
  - to go from source to gdb press i (just like insert mode in vim!!!!!)

Source

GDB

```
591
592    // Ignore spurious request
593    if (size == 0) {
594        dbg_ensures(mm_checkheap(__LINE__));
595        return bp;
596    }
597
598    // Adjust block size to include overhead and to meet alignment requirem
599    asize = round_up(size + dsize, dsize);
600
601    // Search the free list for a fit
602 -> block = find_fit(asize);
603
604    // If no fit is found, request more memory, and then and place the bloc
605    if (block == NULL) {
606        // Always request at least chunksize
607        extendsize = max(asize, chunksize);
608        block = extend_heap(extendsize);
609        // extend_heap returns an error
610        if (block == NULL) {
611            return bp;
612        }
613    }
/afs/andrew.cmu.edu/usr10/jmcamero/private/213_testing/malloclab-s20/mm.c
r -f traces/syn-string-scaled.rep
r -f traces/syn-string-short.rep
r -f traces/syn-string.rep
r -f traces/syn-struct-scaled.rep
r -f traces/syn-struct-short.rep
r -f traces/syn-struct.rep
(gdb) r -f traces/syn-mix-short.rep
Starting program: /afs/andrew.cmu.edu/usr10/jmcamero/private/213_testing/mallocl
ab-s20/mdriver-dbg -f traces/syn-mix-short.rep
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib64/libthread_db.so.1".
Found benchmark throughput 8009 for cpu type Intel(R)Xeon(R)CPUE5520@2.27GHz, be
nchmark regular
Throughput targets: min=4004, max=7208, benchmark=8009

Breakpoint 1, mm_malloc (size=9904) at mm.c:580
(gdb) n
(gdb)
(gdb)
(gdb)
(gdb)
(gdb) _
```

Source Mode

- Benefits
  - see breakpoints in the file
  - reread code while debugging
  - review what is coming
- Usage
  - Very similar to vim!!
  - j - move down a line
  - k - move up a line
  - :### jump to line ###
  - /### search for ###
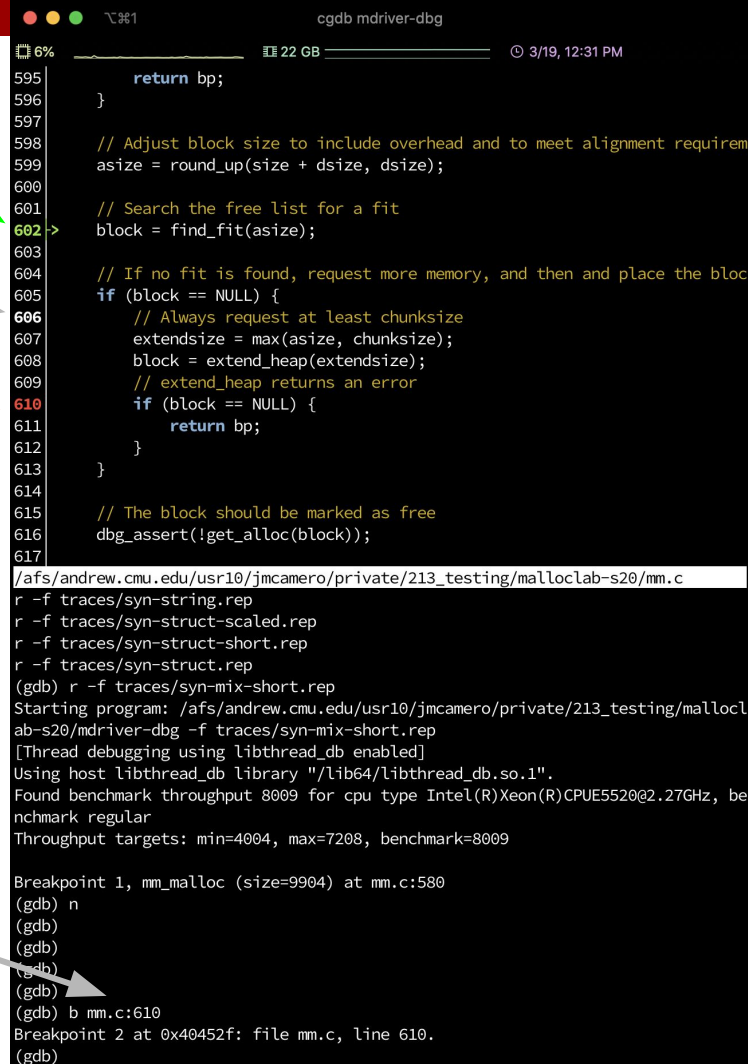 - Most other vim commands!
- Notes
  - Green line number is current line
  - Red line number is breakpoint

Current Line to execute

Viewed line

Next Breakpoint

Breakpoint at mm.c 610

CGDB Misc

- GDB mode functions exactly like normal GDB!
  - All the commands you know and love work the same!
- Shark machines use version 0.6.8
  - Means unfortunately no assembly viewer :(though that is not often needed.
- Website
  - https://cgdb.github.io/
- Documentation
  - https://cgdb.github.io/docs/cgdb.pdf
  - Version 0.7.1

Heap Checker

- int mm_checkheap(int verbose);
- critical for debugging
  - **write this function early!**
  - update it when you change your implementation
  - check all heap invariants, make sure you haven't lost track of any part of your heap
    - check should pass if and only if the heap is truly well-formed
  - should only generate output if a problem is found, to avoid cluttering up your program's output
- meant to be correct, **not** efficient
- call before/after major operations **when the heap should be well-formed**

31

Heap Invariants (**Non-Exhaustive**)

- Block level
  - What are some things which should always be true of every block in the heap?

32

Heap Invariants (**Non-Exhaustive**)

- Block level
  - header and footer match
  - payload area is aligned, size is valid
  - no contiguous free blocks unless you defer coalescing
- List level
  - What are some things which should always be true of every element of a free list?

Heap Invariants (**Non-Exhaustive**)

- Block level
  - header and footer match
  - payload area is aligned, size is valid
  - no contiguous free blocks unless you defer coalescing
- List level
  - next/prev pointers in consecutive free blocks are consistent
  - no allocated blocks in free list, all free blocks are in the free list
  - no cycles in free list unless you use a circular list
  - each segregated list contains only blocks in the appropriate size class
- Heap level
  - What are some things that should be true of the heap as a whole?

34

Heap Invariants (**Non-Exhaustive**)

- Block level

  - header and footer match

  - payload area is aligned, size is valid
  - no contiguous free blocks unless you defer coalescing

- List level

  - next/prev pointers in consecutive free blocks are consistent

  - no allocated blocks in free list, all free blocks are in the free list
  - no cycles in free list unless you use a circular list
  - each segregated list contains only blocks in the appropriate size class

- Heap level

  - all blocks between heap boundaries, correct sentinel blocks (if used)

29

How to Ask for Help

- Be specific about what the problem is, and how to cause it
    - **BAD:** "My program segfaults."
    - **GOOD:** "I ran mdriver in gdb and it says that a segfault occurred due to an invalid next pointer, so I set a watchpoint on the segfaulting next pointer. How can I figure out what happened?"
    - **GOOD:** "My heap checker indicates that my segregated list has a block of the wrong size in it after performing a coalesce(). Why might that be the case?"
    - What sequence of events do you expect around the time of the error? What part of the sequence has already happened?
- Have you written your mm_checkheap function, and is it working?
    - We **WILL** ask to see it!
- Use a rubber duck!

36

If You Get Stuck

- ***Please read the writeup!***
- CS:APP Chapter 9
- View lecture notes and course FAQ at
  http://www.cs.cmu.edu/~213
- Office hours Sunday through Friday 5:30-9:30pm over Zoom. See pinned Piazza post!
- Post a **private** question on Piazza
- Obtain a rubber duck at
 https://tinyurl.com/malloc-f18

APPENDIX

Internal Fragmentation

- Occurs when the *payload* is smaller than the block size

  - due to alignment requirements

  - due to management overhead
  - as the result of a decision to use a larger-than-necessary block

- Depends on the current allocations, i.e. the pattern of *previous* requests

39

Internal Fragmentation

- Due to alignment requirements – the allocator doesn't know how you'll be using the memory, so it has to use the strictest alignment:
  - void *m1 = malloc(13); void *m2 = malloc(11);
  - m1 and m2 both have to be aligned on 8-byte boundaries

- | l | e | n | 1 | p | a | y | l | o | a | d | 1 | l | e | n | 2 | p | a | y | l | o | a | d | 2 |

External Fragmentation

- Occurs when the total free space is sufficient, but no single free block is large enough to satisfy the request
- Depends on the pattern of *future* requests
  - thus difficult to predict, and any measurement is at best an estimate
- Less critical to malloc traces than internal fragmentation

p5 = malloc(4)

free(p1)

p6 = malloc(5)	*Oops!* **Seven bytes available, but not in one chunk....**

C: Pointer Arithmetic

- Adding an integer to a pointer is different from adding two integers
- The value of the integer is always multiplied by the size of the type that the pointer points at
- Example:
  - type_a *ptr = ...;
  - type_a *ptr2 = ptr + a;
- is really computing
  - ptr2 = ptr + (a * sizeof(type_a));
  - i.e.   lea (ptr, a, sizeof(type_a)), ptr2
- Pointer arithmetic on void* is undefined (what's the size of a void?) 42

C: Pointer Arithmetic

- ```
  int *ptr = (int*)0x152130;
  int *ptr2 = ptr + 1;
  ```

- ```
  char *ptr = (char*)0x152130;
  char *ptr2 = ptr + 1;
  ```

- ```
  char *ptr = (char*)0x152130;
  void *ptr2 = ptr + 1;
  ```

- ```
  char *ptr = (char*)0x152130;
  char *p2 = ((char*)(((int*)ptr)+1));
 ```

43

C: Pointer Arithmetic

- `int *ptr = (int*)0x152130;`

`int *ptr2 = ptr + 1;` // ptr2 is 0x152134

- `char *ptr = (char*)0x152130;`
`char *ptr2 = ptr + 1;` // ptr2 is 0x152131

- `char *ptr = (char*)0x152130;`
`void *ptr2 = ptr + 1;` // ptr2 is **still** 0x152131

- `char *ptr = (char*)0x152130;`
`char *p2 = ((char*)(((int*)ptr)+1));` // p2 is 0x152134

44

Dynamic Memory Allocation: Example

p1 = malloc(3)

p2 = malloc(7)

p3 = malloc(5)

free(p2)

p4 = malloc(4)

p5 = malloc(4)

The Memory-Block Information Data Structure

- Requirements:
    - tells us where the blocks are, how big they are, and whether they are free
    - must be able to update the data during calls to malloc and free
    - need to be able to find the **next free block** which is a "good enough fit" for a given payload
    - need to be able to quickly mark a block as free or allocated
    - need to be able to detect when we run out of blocks
        - what do we do in that case?
- The only memory we have is what we're handing out
    - ...but not all of it needs to be payload!   We can use part of it to store the block information.

46

Finding a Free Block

- First Fit
  - search from beginning, use first block that's big enough
  - linear time in total number of blocks
  - can cause small "splinters" at beginning of list
- Next Fit
 - start search from where previous search finished
  - often faster than first fit, but some research suggests worse fragmentation
- Best Fit
  - search entire list, use smallest block that's big enough
  - keeps fragments small (less wasted memory), but slower than first fit

42

Freeing Blocks

- Simplest implementation is just clearing the "allocated" flag
  - but leads to external fragmentation

48

Insertion Policy

- Where do you put a newly-freed block in the free list?
  - LIFO (last-in-first-out) policy
    - add to the beginning of the free list
    - pro: simple and constant time (very fast)
    block->next = freelist; freelist = block;
    - con: studies suggest fragmentation is worse
  - Address-ordered policy
    - insert blocks so that free list blocks are always sorted by address
    addr(prev) < addr(curr) < addr(next)
    - pro: lower fragmentation than LIFO
    - con: requires search

49

C: Pointer Casting

- Notation: (`b*`)`a` "casts" `a` to be of type `b*`
- Casting a pointer doesn't change the bits!
  - `type_a *ptr_a=...; type_b *ptr_b=(type_b*)ptr_a;`
    makes ptr_a and ptr_b contain identical bits
- But it does change the behavior when dereferencing
  - because we *interpret* the bits differently
- Can cast `type_a*` to long/unsigned long and back
  - pointers are really just 64-bit numbers
  - such casts are important for malloclab
  - but be careful – this can easily lead to hard-to-find errors

Cycle Checking: Hare and Tortoise Algorithm

- This algorithm detects cycles in
   linked lists
- Set two pointers, called "hare"
  and "tortoise", to the beginning of
  the list
- During each iteration, move
  "hare" forward by two nodes, "tortoise" by
  one node
  - if "tortoise" reaches the end
    of the list, there is no cycle
  - if "tortoise" equals "hare", the list has a cycle

46

Debugging Tip: Using the Preprocessor

- Use conditional compilation with #if or #ifdef to easily turn debugging code on or off

```
#ifdef DEBUG
#define DBG_PRINTF(...)    fprintf(stderr, __VA_ARGS__)
#define CHECKHEAP(verbose) mm_checkheap(verbose)
#else
#define DBG_PRINTF(...)
#define CHECKHEAP(verbose)
#endif /* DEBUG */
```

```
//  comment line below to disable debug code!
#define DEBUG

void free(void *p) {
    DBG_PRINTF("freeing %p\n", p);
    CHECKHEAP(1);
    ...
}
```

47

Debugging Tip: Using the Preprocessor (contd)

```
#define DEBUG

void free(void *p) {
    DBG_PRINTF("freeing %p\n", p);
    CHECKHEAP(1);

    ...
}
```

preprocessor magic →

```
void free(void *p) {
    fprintf(stderr, "freeing %p\n", p);
    mm_checkheap(1);
    ...
}
```

Replaced with debug code!

```
//  #define DEBUG

void free(void *p) {
    DBG_PRINTF("freeing %p\n", p);
    CHECKHEAP(1);

    ...
}
```

preprocessor magic →

```
void free(void *p) {
    ...
}
```

Debug code gone!

48

Header Reduction

- **Note:** this is completely optional and generally **discouraged** due to its relative difficulty
  - Do **NOT** attempt unless you are satisfied with your implementation as-is

- When to use 8 or 4 byte header? (must support all possible block sizes)
- If 4 byte, how to ensure that payload is aligned?
- Arrange accordingly
- How to coalesce if 4 byte header block is followed by 8 byte header block?
- Store extra information in headers

54