

15-213 Recitation 5

Attack Lab and Stacks

Your TAs

Monday, February 10th, 2020 (15-213, 18-213)

Wednesday, February 12th, 2020 (18-613)

Agenda

- Attack Lab Overview
- Stacks Review
- Activity 1
- Procedure Calling Review
- Activity 2

Learning objectives

By the end of this recitation, we want you to know:

- Stack discipline and calling conventions
- How to perform a simple buffer overflow attack

More discussion in lecture on Tuesday:

Machine-Level Programming V: Advanced Topics

Reminders and Lab Overview

Reminders

- Bomb Lab is due **Tuesday, Feb. 11** at 11pm ET!
- Attack Lab will be released at the same time, and is due the following **Tuesday, Feb. 18**

Attack Lab overview

- Attack programs by crafting buffer overflow attacks that hijack the control flow
- Provide inputs to the rtarget and ctarget programs that cause them to call certain functions
- Unlike in bomblab, the targets don't explode!

Stacks Review

Manipulating the stack

What instructions do we typically use to change the stack pointer, %rsp?

Growing the stack:

Shrinking the stack:

Manipulating the stack`

What instructions do we typically use to change the stack pointer, %rsp?

Growing the stack:

- `sub $0x28, %rsp`
- `push %rbx`
- `callq my_function`

Shrinking the stack:

Manipulating the stack

What instructions do we typically use to change the stack pointer, %rsp?

Growing the stack:

- `sub $0x28, %rsp`
- `push %rbx`
- `callq my_function`

Shrinking the stack:

- `add $0x28, %rsp`
- `pop %rbx`
- `retq`


x86-64 Stack Frames

What kinds of data are stored on the stack?

x86-64 Stack Frames

What kinds of data are stored on the stack?


- Saved registers
- Local variables
- Arguments (7+)
- Saved return address



Which way does the stack grow?

- Up?
- Down?
- Left?
- Right?

Which way does the stack grow?



It depends on how you draw it!

The stack always grows towards **lower addresses** in x86-64.

(Informally, this usually means "down".)

Be aware of this possible ambiguity when reading diagrams.

Drawing memory

Stack diagrams

Carnegie Mellon

Buffer Overflow Stack Example

Before call to gets

```

Stack Frame
for call_echo

00 00 00 00
00 40 06 c3

20 bytes unused

[3] [2] [1] [0] buf ← %rsp
  
```

```

void echo()
{
    char buf[4];
    gets(buf);
    ...
}

echo:
    subq $0x18, %rsp
    movq %rsp, %rdi
    call gets
    ...
  
```

call_echo:

```

    ...
    4006be: callq 4006cf <echo>
    4006c3: add    $0x8,%rsp
    ...
  
```

Addresses are displayed increasing to the left, and then upwards.

Everything else

Carnegie Mellon

Array Example

```

#define ZLEN 5
typedef int zip_dig[ZLEN];

zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
  
```

zip_dig cmu;	1	5	2	1	3
	16	20	24	28	32


zip_dig mit;	0	2	1	3	9
	36	40	44	48	52

zip_dig ucb;	9	4	7	2	0
	56	60	64	68	72

Addresses are displayed increasing to the right, and then downwards.

Endianness

- Describes how integers are represented as bytes.
- Little-endian means that the least-significant 8 bits of an integer are stored at the lowest address.



Endianness

- Describes how integers are represented as bytes.
- Little-endian means that the least-significant 8 bits of an integer are stored at the lowest address.

		A[0]	A[1]	A[2]	A[3]
32-bit integer 0x01020304 "Big end" "Little end"	Little-endian	0x04	0x03	0x02	0x01
	Big-endian	0x01	0x02	0x03	0x04
	Middle-endian	0x02	0x01	0x04	0x03

Activity 1

Part 1: Introduction to solve()

Let's look at `solve()` in
the `src/activity.c` file.

What is it doing?

Is it possible for the
program to call `win()`?

```
void solve(void) {  
    long before = 0xb4;  
    char buf[16];  
    long after = 0xaf;  
  
    Gets(buf);  
  
    if (before == 0x3331323531)  
        win(0x15213);  
  
    if (after == 0x3331323831)  
        win(0x18213);  
}
```

Part 1: The gets() function

```
char *gets(char *s);
```

- `gets()` reads from standard input and writes characters into `s` until it reaches a newline.
- Since it has no information about the **size** of the buffer `s`, its design is fundamentally flawed. **Never use `gets()` yourself!**
- Gets() is a CS:APP wrapper function that checks for errors, and exits if it encounters any.

Part 1: Activity setup

- Split up into groups of 2-3 people
- One person needs a laptop
- Log in to a Shark machine, and type:

```
$ wget https://www.cs.cmu.edu/~213/activities/rec5.tar
$ tar xvf rec5.tar
$ cd rec5
```


- Take a look at the code in `src/activity.c`.

Part 1: Diving into assembly

- Look at the disassembly of solve().
- Try drawing a stack diagram.
 - How large is the stack frame?
 - Where is the saved return address?
 - Where are before, buf, and after?
- **Which variable will be overwritten if we perform a buffer overflow, before or after?**

Part 1: Drawing the stack diagram

```
=> 0x4006b5 <+0>:    sub    $0x38,%rsp
```




Addresses
increase towards
the top of the slide



Part 1: Drawing the stack diagram

```
0x4006b5 <+0>:    sub    $0x38,%rsp  
=> 0x4006b9 <+4>:  movq   $0xb4,0x28(%rsp)
```


Addresses increase towards the top of the slide



Part 1: Drawing the stack diagram

```
0x4006b5 <+0>:    sub    $0x38,%rsp  
0x4006b9 <+4>:    movq   $0xb4,0x28(%rsp)  
=> 0x4006c2 <+13>: movq   $0xaf,0x8(%rsp)
```


Addresses increase towards the top of the slide



Part 1: Drawing the stack diagram

```
0x4006b5 <+0>:    sub    $0x38,%rsp  
0x4006b9 <+4>:    movq    $0xb4,0x28(%rsp)  
0x4006c2 <+13>:   movq    $0xaf,0x8(%rsp)  
0x4006cb <+22>:   lea     0x10(%rsp),%rdi  
=> 0x4006d0 <+27>: callq   0x40073f <Gets>
```


Addresses increase towards the top of the slide



Part 1: Drawing the stack diagram

```
0x4006b5 <+0>:    sub    $0x38,%rsp  
0x4006b9 <+4>:    movq   $0xb4,0x28(%rsp)  
0x4006c2 <+13>:   movq   $0xaf,0x8(%rsp)  
0x4006cb <+22>:   lea    0x10(%rsp),%rdi  
0x4006d0 <+27>:   callq  0x40073f <Gets>  
=> 0x4006d5 <+32>: mov    0x28(%rsp),%rdx
```

Addresses increase towards the top of the slide




Part 1: Comparing with GDB output

Let's compare the stack diagram we drew with the actual values on the stack after Gets() returns.

```
0x4006d0 <+27>:    callq  0x40073f <Gets>
=> 0x4006d5 <+32>:    mov     0x28(%rsp),%rdx
```

```
(gdb) break *0x4006d5
(gdb) run
Starting program: act1
abcdefghijklmnopqrstuvwxyz
(gdb) x/8gx $rsp
(gdb) x/64bx $rsp
```




Part 1: Comparing with GDB output

```
(gdb) x/8gx $rsp
0x602020: 0x0000000000000000      0x00000000000000af
0x602030: 0x686766564636261      0x3837363534333231
0x602040: 0x0000000000000000      0x00000000000000b4
0x602050: 0x0000000000000000      0x000000000400783
```

```
(gdb) x/64bx $rsp
0x602020: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x602028: 0xaf 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x602030: 0x61 0x62 0x63 0x64 0x65 0x66 0x67 0x68
0x602038: 0x31 0x32 0x33 0x34 0x35 0x36 0x37 0x38
0x602040: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x602048: 0xb4 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x602050: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x602058: 0x83 0x07 0x40 0x00 0x00 0x00 0x00 0x00 0x00
```

Addresses
increase towards
bottom of the slide



Part 1: Exploitation

- Try to find an input string that wins 1 cookie!
 - What do we need to overwrite before with if we want to have before == 0x3331323531?
- Constructing an exploit
 - `gets()` stops reading once it sees a newline. In the buffer, it replaces the newline with a null terminator.
 - `gets()` does **not** stop reading at a null terminator.

Part 1: Recap

- Buffer overflows can **overwrite** parts of the stack frame, including other local variables
- Stack frames may include **padding**, so looking at the assembly is crucial to drawing a correct diagram
- GDB prints output starting at the **lowest** address, whereas our stack diagrams start at the **highest**

Procedure Calling Review

Call and return instructions

Which registers do callq and retq change?

%rax

%rdi

%rsi

%rdx

%rcx

%r8

%r9

%r10

%r11

%rbx

%r12

%r13

%r14



%rbp

%rsp

%rip


Call and return instructions

Which registers do callq and retq change?



Stack/Procedure Review

```
000000000400540 <multstore>:  
    .  
    .  
=>400544: callq  400550 <mult2>  
400549: mov     %rax, (%rbx)  
    .  
    .
```




%rsp **0x120**

%rip **0x400544**

```
000000000400550 <mult2>:  
400550: mov     %rdi,%rax  
    .  
    .  
400557: retq
```

Stack/Procedure Review

```
000000000400540 <multstore>:  
    .  
    .  
=>400544: callq  400550 <mult2>  
400549: mov     %rax, (%rbx)  
    .  
    .
```



%rsp 0x120

%rip 0x400544


```
000000000400550 <mult2>:  
400550: mov     %rdi,%rax  
    .  
    .  
400557: retq
```

What happens next?

Stack/Procedure Review

```
000000000400540 <multstore>:  
    .  
    .  
400544: callq  400550 <mult2>  
400549: mov     %rax, (%rbx)  
    .  
    .
```


```
000000000400550 <mult2>:  
=>400550:  mov     %rdi,%rax  
    .  
    .  
400557:  retq
```



Stack/Procedure Review

```
000000000400540 <multstore>:  
    .  
    .  
400544: callq  400550 <mult2>  
400549: mov     %rax, (%rbx)  
    .  
    .
```


```
000000000400550 <mult2>:  
    400550: mov     %rdi,%rax  
    .  
    .  
=>400557: retq
```



Stack/Procedure Review

```
000000000400540 <multstore>:  
    .  
    .  
    .  
    400544: callq  400550 <mult2>  
=>400549: mov     %rax, (%rbx)  
    .  
    .
```


```
000000000400550 <mult2>:  
    400550: mov     %rdi,%rax  
    .  
    .  
    400557: retq
```



Let's Rewind...

```
000000000400540 <multstore>:  
    .  
    .  
400544: callq  400550 <mult2>  
400549: mov     %rax, (%rbx)  
    .  
    .
```

```
000000000400550 <mult2>:  
400550: mov     %rdi,%rax  
    • ????  
    • ????  
=>400557: retq
```



What if we mess up the return address?

Activity 2

Part 2: Exploitation

- Hijacking control flow
 - Is it possible to overwrite after? If not, what parts of the stack frame *can* we overwrite?
 - Is there anywhere we could jump to call `win(0x18213)`?
- Constructing an exploit

inputs/input2.txt

```
48 65 6c 6c 6f 20 31 35  
32 31 33 21 # comment
```



`make`
(runs hex2raw)

inputs/input2.bin

```
Hello 15213!
```

Part 2: Recap

- `retq` always jumps to the **saved return address**, which it pops off the stack (at `rsp`).
- **Overwriting** the saved return address on the stack allows us to "fool" `retq`, and transfer control to an arbitrary instruction.

Attack Lab Tools

- ```
$ gcc -c test.s
$ objdump -d test.o
```

Compiles the assembly code in test.s, then shows the disassembled instructions along with the actual bytes.
- ```
$ ./hex2raw < exploit.txt > exploit.bin
```

Convert hex codes into raw binary strings to pass to targets.
- ```
(gdb) display /12gx $rsp
(gdb) display /2i $rip
```

Displays 12 elements on the stack and the next 2 instructions to run  
GDB is also useful to for tracing to see if an exploit is working.

# If you get stuck

- **Please read the writeup carefully.** Not everything will make sense on the first read-through.
- Other resources you can make use of:
  - CS:APP Chapter 3
  - Lecture slides and videos
  - x86-64 and GDB cheat sheets under [Resources](#)