

Recitation 10: Malloc Lab

Your TAs

Monday, March 16th, 2020

Course Updates


- **Office hours**
 - Reference the post and video on piazza
- **Future recitations**
 - Join the same way you joined this one!
- **More responsive on piazza**
- **Lecture plans**
 - Tuesday, 1:30 PM: Discussion of courses changes with Saugata
 - Flipped lectures starting Thursday
- **We will give you more details via piazza when we have updates!**

Administrivia

- Malloc traces due tomorrow Tuesday, March 17!
- Malloc checkpoint due Tuesday, March 24! yeeT
- Malloc final due Tuesday, March 31! yooT
- Malloc Bootcamp Thursday, March 19 (@ 6pm)!

Traces Assignment

- Due tomorrow! (March 17)
- *Read the writeup!* (useful things like trace file format)
- Write 3 test cases that trigger different coalesce cases for merging newly freed blocks



- Understand how coalesce works, since you'll have to implement it when you write your own malloc :)

Checkpoint Submission

■ Style Grading

- We will grade your checkheap with your checkpoint submission!

■ Things to Remember:

- Document checkheap
- See writeup for what to include in checkheap

Git Reminders

- **Style grades for CacheLab have been released! Points were also deducted for style on Git usage**
 - Please use detailed commit messages – things like “DONE” or “did a thing” aren’t enough
 - You should be committing often as you work on your code
 - Especially for malloc: `git diff` can show what you changed since your last working commit
 - Also allows you to restore your hard work in case your file gets deleted accidentally...
- **Commit early, commit often** 
- **Remember to git push your commits, or else we can't grade them :(**


Outline

- Concept
- How to choose blocks
- Metadata
- Debugging / GDB Exercises

What is malloc?

- A function to allocate memory during runtime (dynamic memory allocation).
 - More useful when the size or number of allocations is unknown until runtime (e.g., data structures)

- The heap is a segment of memory addresses reserved almost exclusively for malloc to use.
 - Your code directly manipulates the bytes of memory in this section.




Concept

- Overall, malloc does three things:
 1. Organizes all blocks and stores information about them in a structured way.
 2. Uses the structure made to choose an appropriate location to allocate new memory.
 3. Updates the structure when the user frees a block of memory.

This process occurs even for a complicated algorithm like segregated lists.

Concept (Implicit list)


1. Connects and organizes all blocks and stores information about them in a structured way, typically implemented as a singly linked list




Concept (Implicit list)

2. Uses the structure made to choose an appropriate location to allocate new memory.

p1 = malloc(3)



p2 = malloc(7)




p3 = malloc(5)




Concept (Implicit list)

- Updates the structure when the user frees a block of memory.


p1 = malloc(3)




p2 = malloc(7)



p3 = malloc(5)



free(p2)




Concept (Implicit list)

- Updates the structure when the user frees a block of memory.


p1 = malloc(3)




p2 = malloc(7)






p3 = malloc(5)




free(p2)




block to
be freed




Coalesce: Case 1



Result:




Coalesce: Case 2




Result:




Coalesce: Case 3




Result:



Coalesce: Case 4



Result:



Goals

- Run as fast as possible
- Waste as little memory as possible
- Seemingly conflicting goals, but with ~~the library malloc~~ all cleverness you can do very well in both areas!
- The simplest implementation is the implicit list.
mm.c uses this method.
 - Unfortunately...

```
[dalud@angelshark:~/.../15213/sl7/malloclabcheckpoint-handout] $ ./mdriver -p
Found benchmark throughput 13090 for cpu type Intel(R)Xeon(R)CPUE5520@2.27GHz, benchmark checkpoint
Throughput targets: min=2618, max=11781, benchmark=13090
.....
Results for mm malloc:
  valid   util     ops    msecs   Kops   trace
  yes    78.4%      20     0.002   9632 ./traces/syn-array-short.rep
  yes    13.4%      20     0.001  25777 ./traces/syn-struct-short.rep
  yes    15.2%      20     0.001  24783 ./traces/syn-string-short.rep
  yes    73.1%      20     0.001  19277 ./traces/syn-mix-short.rep
  yes    16.0%      36     0.001  31192 ./traces/ngram-foxl.rep
  yes    73.6%     757     0.145   5237 ./traces/syn-mix-realloc.rep
* yes   62.0%     5748     3.925   1464 ./traces/bdd-aa4.rep
* yes   58.3%   87830   1682.766    52 ./traces/bdd-aa32.rep
* yes   58.0%   41080   410.385   100 ./traces/bdd-ma4.rep
* yes   58.1%  115380   4636.711    25 ./traces/bdd-nq7.rep
* yes   56.6%   20547    26.677   770 ./traces/cbit-abs.rep
* yes   55.8%   95276    675.303   141 ./traces/cbit-parity.rep
* yes   58.0%   89623    611.511   147 ./traces/cbit-satadd.rep
* yes   49.6%   50583    185.382   273 ./traces/cbit-xyz.rep
* yes   40.6%   32540    76.919   423 ./traces/ngram-gulliver1.rep
* yes   42.4%  127912   1284.959   100 ./traces/ngram-gulliver2.rep
* yes   39.4%   67012    338.591   198 ./traces/ngram-moby1.rep
* yes   38.6%   94828    701.305   135 ./traces/ngram-shakel.rep
* yes   90.9%  80000   1455.891    55 ./traces/syn-array.rep
* yes   88.0%  80000    915.167    87 ./traces/syn-mix.rep
* yes   74.3%  80000    914.366    87 ./traces/syn-string.rep
* yes   75.2%  80000    812.748    98 ./traces/syn-struct.rep
16 16   59.1% 1148359  14732.604    78

Average utilization = 59.1%. Average throughput = 78 Kops/sec
Checkpoint Perf index = 20.0 (util) + 0.0 (thru) = 20.0/100
```

This is pretty
slow... most
explicit list
implementations
get above 2000
Kops/sec

Allocation methods in a nutshell

- **Implicit list:** a list is implicitly formed by jumping between blocks, using knowledge about their sizes.



- **Explicit list:** Free blocks explicitly point to other blocks, like in a linked list.

- Understanding explicit lists requires understanding implicit lists



- **Segregated list:** Multiple linked lists, each containing blocks in a certain range of sizes.

- Understanding segregated lists requires understanding explicit lists



Choices

■ What kind of implementation to use?

- Implicit list, explicit list, segregated lists, binary tree methods, etc.
- You can use specialized strategies depending on the size of allocations
- Adaptive algorithms are fine, though not necessary to get 100%.
 - Don't hard-code for individual trace files - you'll get no credit/code deductions!

■ What fit algorithm to use?

- Best fit: choose the smallest block that is big enough to fit the requested allocation size
- First fit / next fit: search linearly starting from some location, and pick the first block that fits.
- Which is faster? Which uses less memory?
- “Good enough” fit: a blend between the two

■ This lab has many more ways to get an A+ than, say, Cache Lab Part 2

Finding a Best Block

- Suppose you have implemented the explicit list approach
 - You were using best fit with explicit lists
- You experiment with using segregated lists instead.
Still using best fits.
 - Will your memory utilization score improve?

Note: you don't have to implement seglists and run mdriver to answer this. That's, uh, hard to do within one recitation session.

- What other advantages does segregated lists provide?
- Losing memory because of the way you choose your free blocks is called external fragmentation.

Metadata

- All blocks need to store some data about themselves in order for malloc to keep track of them (e.g. headers)
 - This takes memory too...
 - Losing memory for this reason is called internal fragmentation.
- What data might a block need?
 - Does it depend on the malloc implementation you use?
 - Is it different between free and allocated blocks?
- Can we use the extra space in free blocks?
 - Or do we have to leave the space alone?
- How can we overlap two different types of data at the same location?

In a perfect world...

Setting up the blocks, metadata, lists... etc (500 LoC)

- + Finding and allocating the right blocks (500 LoC)**
- + Updating your heap structure when you free (500 LoC) =**


```
[dalud@angelshark:~/.../15213/s17/malloclabcheckpoint-handout] $ ./mdriver
Found benchmark throughput 13056 for cpu type Intel(R) Xeon(R) CPU E5520 @ 2.27G
Throughput targets: min=6528, max=11750, benchmark=13056
.....
Results for mm malloc:
  valid    util      ops     msecs      Kops   trace
    yes    78.1%      20      0.004     5595 ./traces/syn-array-short.rep
    yes    3.2%       20      0.004     5273 ./traces/syn-struct-short.rep
  * yes   96.0%    80000     17.176     4658 ./traces/syn-array.rep
  * yes   93.2%    80000      6.154    12999 ./traces/syn-mix.rep
  * yes   86.4%    80000      3.717    21521 ./traces/syn-string.rep
  * yes   85.6%    80000      3.649    21924 ./traces/syn-struct.rep
16 16    74.2% 1148359     55.949    20525

Average utilization = 74.2%. Average throughput = 20525 Kops/sec
Perf index = 60.0 (util) + 40.0 (thru) = 100.0/100
```

In reality...

- Setting up the blocks, metadata, lists... etc (500 LoC)
- + Finding and allocating the right blocks (500 LoC)
- + Updating your heap structure when you free (500 LoC)
- + One bug, somewhere lost in those 1500 LoC =

```
[dalud@angelshark:~/.../15213/s17/malloclabcheckpoint-handout] $ ./mdriver
Found benchmark throughput 13056 for cpu type Intel(R) Xeon(R) CPU E5520 @ 2.27GHz
Throughput targets: min=6528, max=11750, benchmark=13056
....Segmentation fault
[dalud@angelshark:~/.../15213/s17/malloclabcheckpoint-handout] $ █
```



Common errors you might see

■ Garbled bytes

- Problem: overwriting data in an allocated block
- Solution: ~~remembering data lab and the good ol' days~~ finding where you're overwriting by stepping through with gdb

■ Overlapping payloads

- Problem: having unique blocks whose payloads overlap in memory
- Solution: ~~literally print debugging everywhere~~ finding where you're overlapping by stepping through with gdb

■ Segmentation fault

- Problem: accessing invalid memory
- Solution: ~~crying a little~~ finding where you're accessing invalid memory by stepping through with gdb

■ Try running \$ make

- If you look closely, our code compiles your malloc implementation with the `-O3` flag.
- This is an optimization flag. `-O3` makes your code run as efficiently as the compiler can manage, but also makes it horrible for debugging (almost everything is “optimized out”).

```
[dalud@angelshark:~/.../15213/s17/rec11] $ make
gcc -Wall -Wextra -Werror -O3 -g -DDRIVER -Wno-unused-function -Wno-u
./macro-check.pl -f mm.c
clang -Wall -Wextra -Werror -O3 -g -DDRIVER -Wno-unused-function -Wno-u
gcc -Wall -Wextra -Werror -O3 -g -DDRIVER -Wno-unused-function -Wno-u
```

```
(gdb) print block
$3 = <optimized out>
(gdb) print asize
$4 = <optimized out>
```

- For malloclab, we've provided you a driver, `mdriver-dbg`, that not only enables debugging macros, but compiles your code with `-O0`. This allows more useful information to be displayed in GDB

Debugging Strategies

- **Write a heap checker!**

- Checks the invariants of your heap to make sure everything is well-formed
- If you write detailed error messages, you can see exactly why your heap is incorrectly formed

- **Use assertions in your functions!**

- 122 style contracts can also help you catch where things go amiss
- Gives more information than a segfault
- Import

- **Use a debugger!**

Debugging Guidelines

If you have this problem...

Ran into segfault



You might want to...

Locate a segfault

- run
- <>
- backtrace
- list

Trace results don't match yours



Don't know what trace output
should be



Reproduce results of a trace

- Run with gdb
- gdb args

What's better than printf? Using GDB

- Use GDB to determine where segfaults happen!
- **gdb mdriver** will open the malloc driver in gdb
 - Type run and your program will run until it hits the segfault!
- **step/next** - (abbrev. **s/n**) step to the next line of code
 - **next** steps over function calls
- **finish** - continue execution until end of current function, then break
- **print <expr>** - (abbrev. **p**) Prints **any C-like expression** (including results of function calls!)
 - Consider writing a heap printing function to use in GDB!
- **x <expr>** - Evaluate <expr> to obtain address, then examine memory at that address
 - **x /a <expr>** - formats as address
 - See **help p** and **help x** for information about more formats

Debugging mdriver

- **(gdb) x /gx block**
 - Shows the memory contents within the block
 - In particular, look for the header.
- **(gdb) print *block**
 - Alternative: **(gdb) print *(block_t *) <address>**
 - Shows struct contents

Using GDB - Fun with frames

- **backtrace** - (abbrev. **bt**) print call stack up until current function
 - **backtrace full** - (abbrev. **bt full**) print local variables in each frame

```
(gdb) backtrace
#0  find_fit (...)

#1  mm_malloc (...)

#2  0x0000000000403352 in eval_mm_valid
(...) #3 run_tests (...)

#4  0x0000000000403c39 in main (...)
```

- **frame 1** - (abbrev. **f 1**) switch to mm_malloc's stack frame
 - Good for inspecting local variables of calling functions

Using GDB - Setting breakpoints/watchpoints

- **break mm_checkheap** - (abbrev. **b**) break on “mm_checkheap()”
 - **b mm.c:25** - break on line 25 of file “mm.c” - **very useful!**
- **b find_fit if size == 24** - break on function “find_fit()” if the local variable “size” is equal to 24 - **“conditional breakpoint”**
- **watch heap_listp** - (abbrev. **w**) break if value of “heap_listp” changes - **“watchpoint”**
- **w block == 0x80000010** - break if “block” is equal to this value
- **w *0x15213** - watch for changes at memory location 0x15213
 - Can be *very slow*
- **rwatch <thing>** - stop on reading a memory location
- **awatch <thing>** - stop on *any* memory access

Heap consistency checker

- mm-2.c activates debug mode, and so mm_checkheap runs at the beginning and end of many of its functions.

```
106 /*  
107  * If DEBUG is defined, enable printing on dbg_printf and contracts.  
108  * Debugging macros, with names beginning "dbg_" are allowed.  
109  * You may not define any other macros having arguments.  
110 */  
111 #define DEBUG // uncomment this line to enable debugging  
112  
113 #ifdef DEBUG  
114 /* When debugging is enabled, these form aliases to useful functions */  
115 #define dbg_printf( ) printf( VA_ARGS )
```

Heap Checker

- `int mm_checkheap(int verbose);`
- critical for debugging
 - **write this function early!**
 - update it when you change your implementation
 - check all heap invariants, make sure you haven't lost track of any part of your heap
 - check should pass if and only if the heap is truly well-formed
 - should only generate output if a problem is found, to avoid cluttering up your program's output
- meant to be correct, **not efficient**
- call before/after major operations **when the heap should be well-formed**

Heap Invariants (Non-Exhaustive)

- Block level
 - What are some things which should always be true of every block in the heap?

Heap Invariants (Non-Exhaustive)

- Block level
 - header and footer match
 - payload area is aligned, size is valid
 - no contiguous free blocks unless you defer coalescing
- List level
 - What are some things which should always be true of every element of a free list?

Heap Invariants (Non-Exhaustive)

- Block level
 - header and footer match
 - payload area is aligned, size is valid
 - no contiguous free blocks unless you defer coalescing
- List level
 - next/prev pointers in consecutive free blocks are consistent
 - no allocated blocks in free list, all free blocks are in the free list
 - no cycles in free list unless you use a circular list
 - each segregated list contains only blocks in the appropriate size class
- Heap level
 - What are some things that should be true of the heap as a whole?

Heap Invariants (Non-Exhaustive)

- Block level
 - header and footer match
 - payload area is aligned, size is valid
 - no contiguous free blocks unless you defer coalescing
- List level
 - next/prev pointers in consecutive free blocks are consistent
 - no allocated blocks in free list, all free blocks are in the free list
 - no cycles in free list unless you use a circular list
 - each segregated list contains only blocks in the appropriate size class
- Heap level
 - all blocks between heap boundaries, correct sentinel blocks (if used)

Strategy - Suggested Plan for Completing Malloc

0. *Start writing your checkheap!*

- 1. Get an explicit list implementation to work with proper coalescing and splitting**
- 3. Get to a segregated list implementation to improve utilization**
- 4. Work on optimizations (each has its own challenges!)**
 - Remove footers**
 - Decrease minimum block size**
 - Reduce header sizes**

Strategy - Suggested Plan for Completing Malloc

- 0. Start writing your *checkheap!*** *Keep writing your checkheap!*
- 1. Get an explicit list implementation to work with proper coalescing and splitting** *Keep writing your checkheap!*
- 3. Get to a segregated list implementation to improve utilization**
Keep writing your checkheap!
- 4. Work on optimizations (each has its own challenges!)**
 - Remove footers *Keep writing your checkheap!*
 - Decrease minimum block size
 - Reduce header sizes

MallocLab Checkpoint

- Due next Tuesday!
- Checkpoint should take a bit less than half of the time you spend overall on the lab.
- Read the write-up. Slowly. Carefully.
- Use GDB - watch, backtrace
- Ask us for debugging help
 - Only after you implement mm_checkheap though! You gotta learn how to understand your own code - help us help you!



please write checkheap
or we will scream

Appendix: Advanced GDB Usage

- **backtrace**: Shows the call stack
- **up/down**: Lets you go up/down one level in the call stack
- **frame**: Lets you go to one of the levels in the call stack
- **list**: Shows source code
- **print <expression>**:
 - Runs any valid C command, even something with side effects like `mm_malloc(10)` or `mm_checkheap(1337)`
- **watch <expression>**:
 - Breaks when the value of the expression changes
- **break <function / line> if <expression>**:
 - Only stops execution when the expression holds true
- **Ctrl-X Ctrl-A or cgdb for visualization**