

Recitation 13: ProxyLab Part 1

Instructor: TA(s)

Outline

- Feedback
- Proxies
- Networking
- PxyDrive Demo

Midsemester + TA Feedback Form

- Course Feedback Form (10 min)
- TA Feedback Form (5 min each)
- Both links are in a pinned piazza post

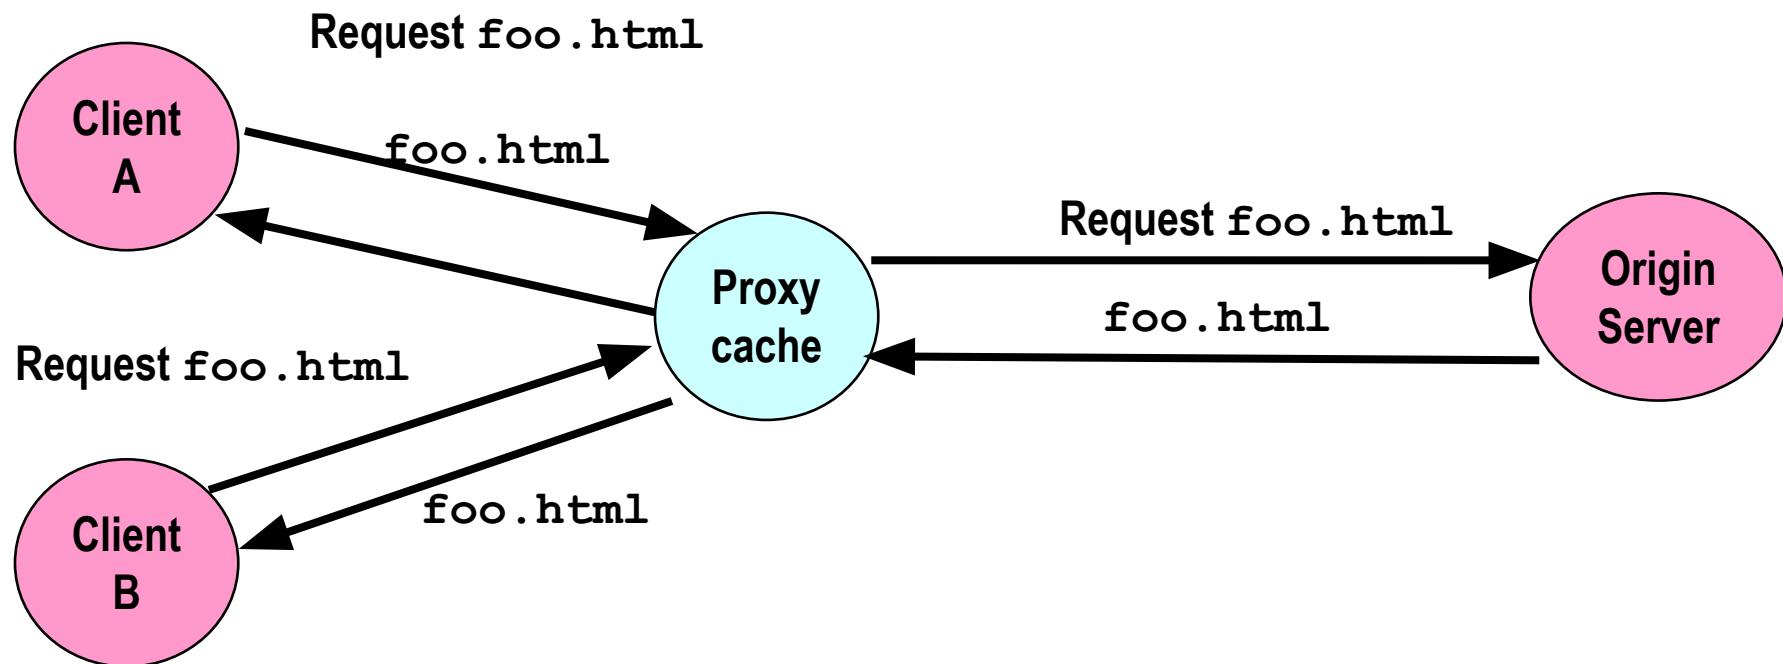
**Please give us feedback so we can improve the course
and be better TAs! :)**

Proxy Lab

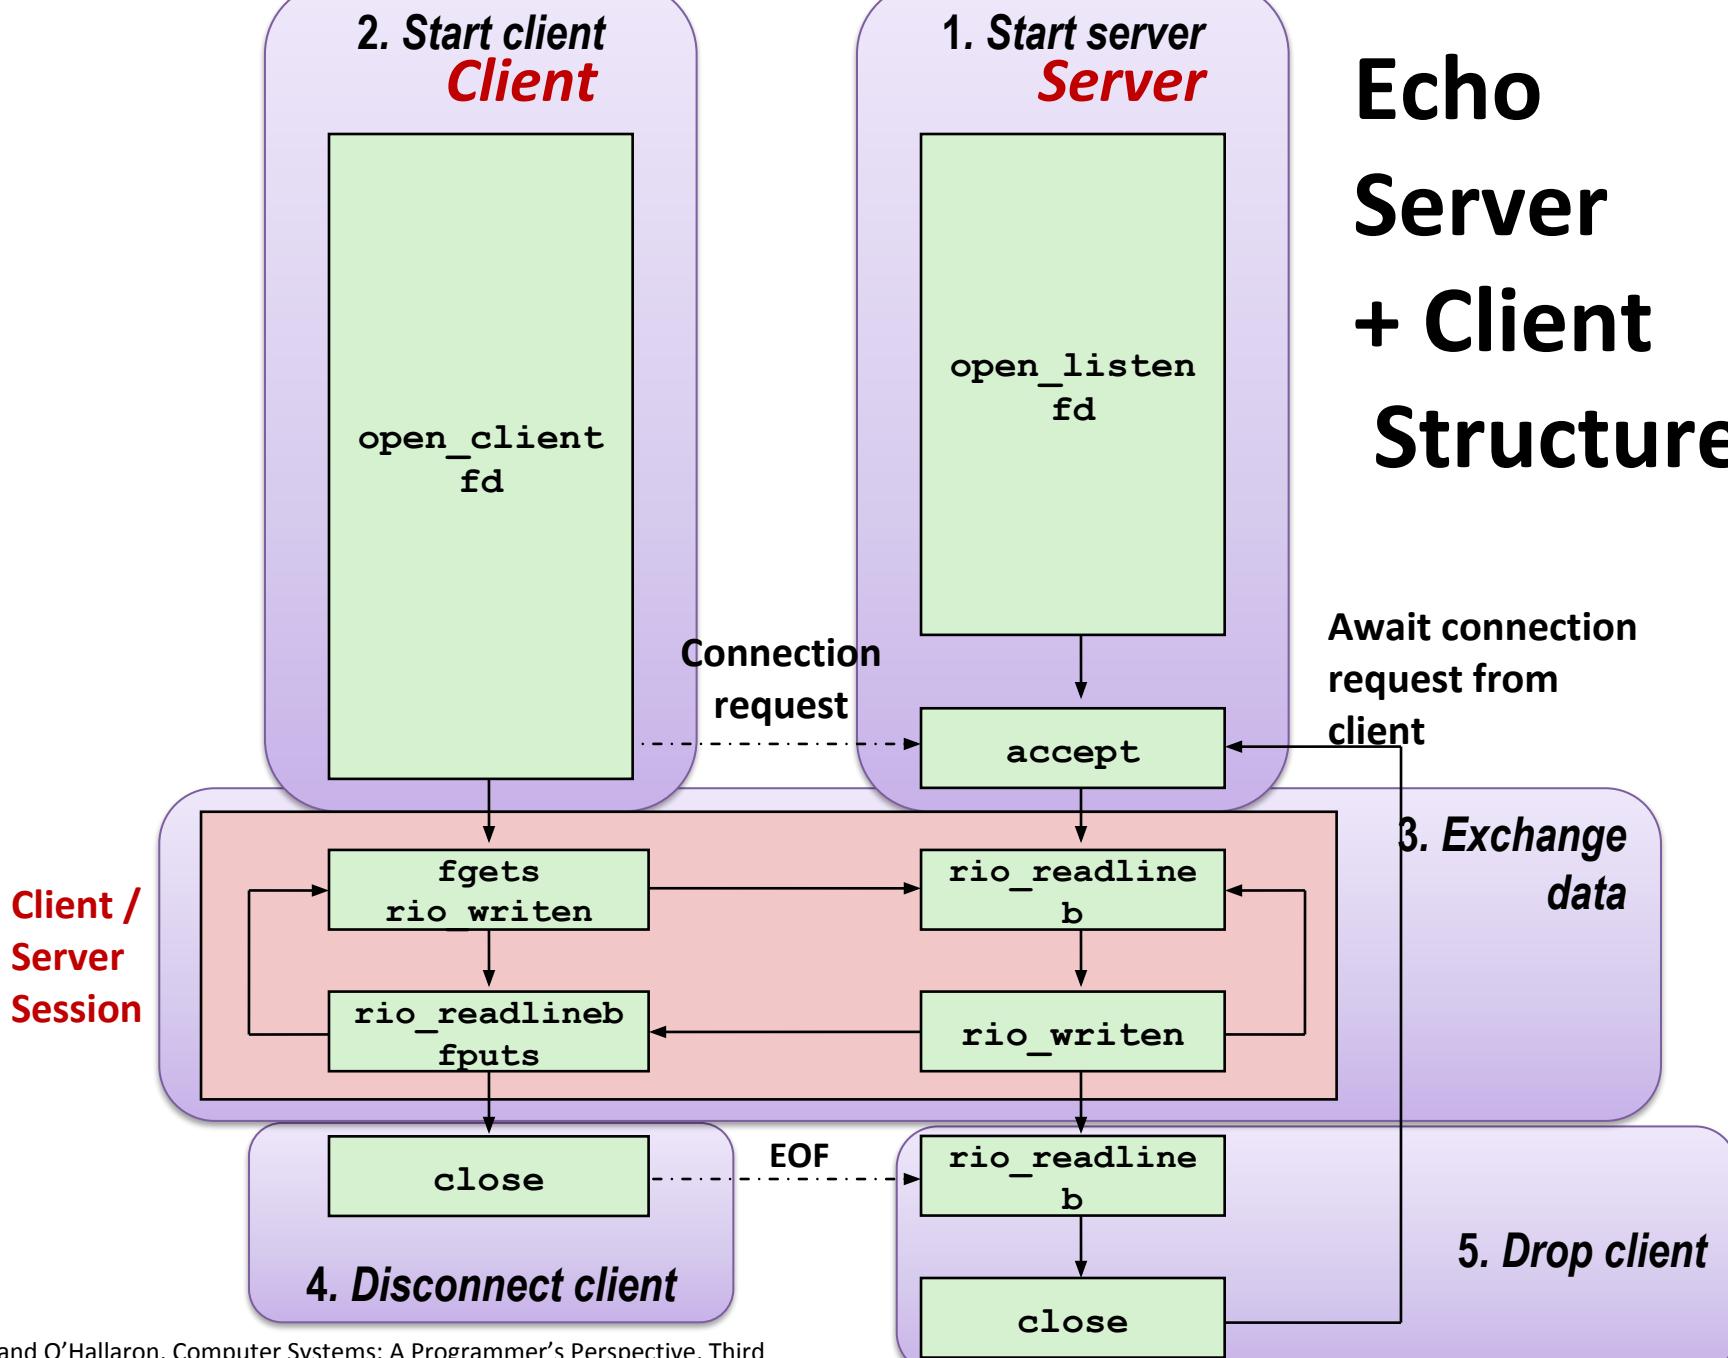
- **Checkpoint is worth 2%, due Thursday, April 23rd**
- **Final is worth 6%, due Thursday, April 30th**
- **Current situation w/ grace / late days (subject to change):**
 - 1 grace / late day allowed for both checkpoint and final
- **You are submitting an entire project**
 - Modify the makefile
 - Split source file into separate pieces
- **Submit regularly to verify proxy builds on Autolab**
- **Your proxy is a server, it should not crash!**

Why Proxies?

- Proxies are both clients and servers
- Can perform useful functions as requests and responses pass by
 - Examples: Caching, logging, anonymization, filtering, transcoding



Echo Server + Client Structure



Transferring HTTP Data

If something requests a file from a web server, how does it know that the transfer is complete?

- A) It reads a NULL byte.
- B) The connection closes.
- C) It reads a blank line.
- D) The HTTP header specifies the number of bytes to receive.
- E) The reading function receives EOF.



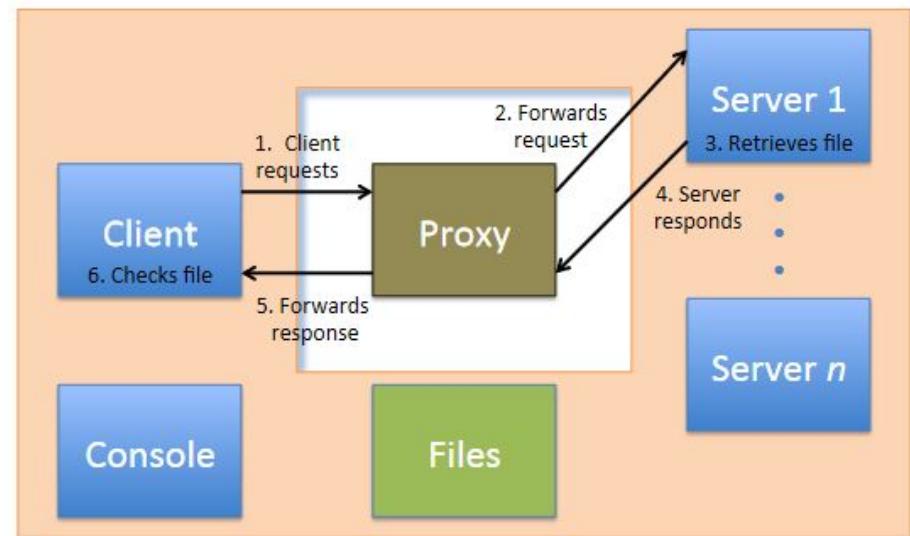
Introducing PxyDRIVE¹

- A REPL for testing your proxy implementation
 - We also grade using this
- Typical pre-f18 proxy debugging experience:
 - Open up three terminals:
for Tiny server, **gdb proxy** and curl
 - Can make multiple requests, but need more terminals
for multiple instances of the Tiny server
 - If the data is corrupted, need to manually inspect lines
of gibberish binary data to check error
- Not anymore with PxyDRIVE!

¹ Not typing PxyDRIVE in small-caps is a style violation.

Introducing PxyDRIVE

- General workflow
 - Generate text and binary data to test your proxy with
 - Create (multiple) servers
 - Make **transactions**
 - Trace transactions to inspect headers and response data
- Transaction



Some practice

- Get the tarball
- \$ wget
<http://www.cs.cmu.edu/~213/activities/proxy-recitation13.tar>
- \$ tar -xvf proxy-recitation13.tar
- \$ cd pxydrive-tutorial

Trying out PxyDRIVE

- It's a REPL: the user can run commands
- `$./pxy/pxydrive.py`
 - Just starts PxyDRIVE
 - Try entering commands:
 - `>help`
 - `>help help help help help...`
 - `>quit`
- `$./pxy/pxydrive.py -p ./proxy-ref`
 - Starts PxyDRIVE and specifies a proxy to run
 - **Proxy set up at <someshark>:30104**
 - Picks the right port and starts the proxy
 - `./proxy-ref` is the reference proxy

PxyDRIVE Tutorial 1

- Introducing basic procedures:
generate data, create server, fetch / request file from server,
trace transaction
- Open `s01-basic-fetch.cmd`

PxyDRIVE Tutorial 1

- **>generate data1.txt 1K**
 - Generates a 1K text file called *data1.txt*
- **>serve s1**
 - Launches a server called *s1*
- **>fetch f1 data1.txt s1**
 - Fetches *data1.txt* from server *s1*, in a transaction called *f1*
- **>wait ***
 - Waits for all transactions to finish
 - Needed in the trace, not in the command-line
- **>trace f1**
 - Traces the transaction *f1*
- **>check f1**
 - Checks the transaction *f1*

PxyDRIVE Tutorial 1

- Run trace with -f option:
- \$./pxy/pxydrive.py -p ./proxy-ref
-f s01-basic-fetch.cmd

Look at the trace of the transaction!

- Identify:
 - GET command
 - Host header
 - Other headers
 - Request from client to proxy
 - Request from proxy to server
 - Response by server to proxy
 - Response by proxy to client

PxyDRIVE Tutorial 1

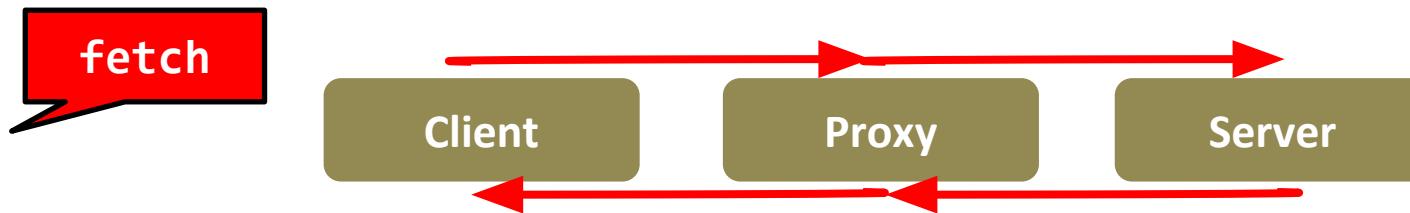
- Run a different trace
- `$./pxy/pxydrive.py -p ./proxy-ref
-f s02-basic-request.cmd`
- You should get a different output from the first trace
- Why? Let's look at this trace...

PxyDrive Tutorial 1

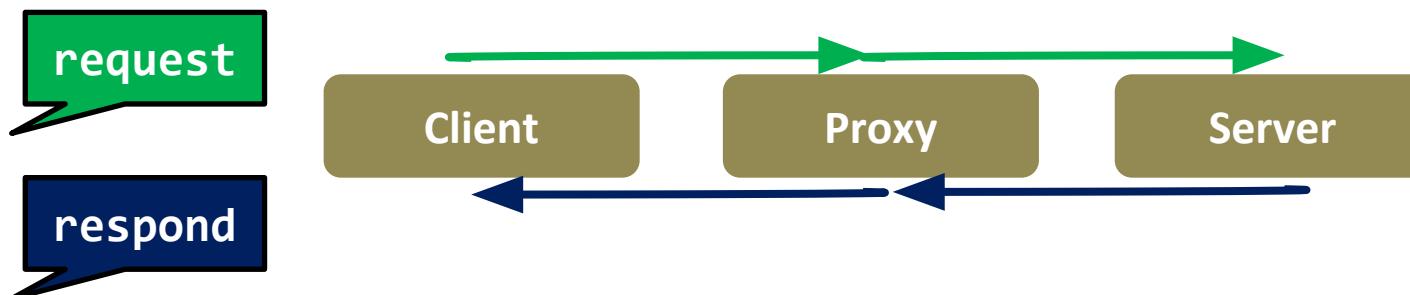
- **>generate data1.txt 1K**
- **>serve s1**
- **>request r1 data1.txt s1**
 - Requests *data1.txt* from server *s1*, in a transaction called *r1*
- **>wait ***
- **>trace r1**
- **>respond r1**
 - Allow server to respond to the transaction *r1*
- **>wait ***
- **>trace r1**
- **>check r1**
 - Checks the transaction *r1*

PxyDRIVE Tutorial 1

- The fetch command makes the server immediately respond to a request.
- All steps of a transaction are complete after a fetch.



- The request command does not complete a transaction.
- A request needs a respond to complete its transaction.



PxyDRIVE Tutorial 2

- Debugging a proxy that clobbers responses
- Run the same trace but with a faulty proxy
- `$./pxy/pxydrive.py -f s01-basic-fetch.cmd
-p ./proxy-corrupt`

What went wrong?

```
Response status: ok
  Source file in ./source_files/random/data1.txt
Request status: error (Mismatch between source file ./source_files/random/data1
.txt and response file ./response_files/f1-data1.txt starting at position 447: 'F' (hex 0x46) ≠ 'G' (hex 0x47))
  Result file in ./response_files/f1-data1.txt
>#
># Make sure it was retrieved properly
>check f1
ERROR: Request f1 generated status 'error'.  Expecting 'ok' (Mismatch between so
urce file ./source_files/random/data1.txt and response file ./response_files/f1-
data1.txt starting at position 447: 'F' (hex 0x46) ≠ 'G' (hex 0x47))
>quit
ERROR COUNT = 1
-bash-4.2$ _
```

PxyDRIVE Tutorial 3

- Debugging a proxy that clobbers headers
- Run the same trace but with another faulty proxy
- ```
$./pxy/pxydrive.py -f s01-basic-fetch.cmd
-p ./proxy-strip -S 3
```
- **-S specifies strictness level**

# What went wrong?

```
Response status: bad_request (Missing Request-ID header)
 Source file in ./source_files/random/data1.txt
Request status: bad_request (bad request)
 Result file in ./response_files/f1-status.html
>#
># Make sure it was retrieved properly
>check f1
ERROR: Request f1 generated status 'bad_request'. Expecting 'ok' (Bad request)
>quit
ERROR COUNT = 1
-bash-4.2$ _
```

# PxyDRIVE Tutorial 4

- Debugging a proxy that crashes
- Run the same trace but with yet another faulty proxy
- `$ ./pxy/pxydrive.py -f s03-overrun.cmd  
-p ./proxy-overrun`
- Is the error message helpful?

# PxyDRIVE Tutorial 4

- We resort to multi-window debugging
- Set up another window and run GDB in one:
  
- `$ gdb ./proxy-overrun`
- `(gdb) run <port>`
  
- In the other window, run PxyDRIVE:
  - `$ ./pxy/pxydrive.py -P localhost:<port> -f s03-overrun.cmd`
  - -P specifies the host and port the proxy is running on

`./port-for-user.pl`  
Run this to get your  
unique port!

# Reminders

- **Read the writeup**
- **One grace / late day for both checkpoint and final**
- **So you really have to start early**
  - Come to office hours this week, before it gets crowded!
- **Work incrementally and take breaks**
- **Simpler tests should be completed in the first week!**



# So you wanna TA for 213?



## What qualifications are we looking for?

- Decent class performance, but also critical thinking skills
- Like computer systems + want to help others like systems!
- Have a reasonable ability to gauge your schedule + responsibilities
- Leadership potential! Take initiative, we love to see it 😊
- Ability to tell students:
  - “Did you write your heap checker”
  - “Run backtrace for me”
  - rinse and repeat, it’s mouthwash baby

# Appendix on echoserver / client Echoserver, echoclient

# Echo Demo

- See the instructions written in the telnet results to set up the echo server. Get someone nearby to connect using the echo client.
- What does echoserver output? (Sample output:)

```
$./echoserver 10101
Accepted connection from hammerheadshark.ics.cs.cmu.edu:46422
hammerheadshark.ics.cs.cmu.edu:46422 sent 6 bytes
Disconnected from hammerheadshark.ics.cs.cmu.edu:46422
```

Server  
listening port

Client  
host

Client  
port

# Echo Demo

## ■ Look at `echoclient.c`

- Opens a connection to the server
- Reads/writes from the server

## ■ Look at `echoserver` output

- Why is the printed client port different from the server's listening port?
- Server opens **one “listening” port**
  - Incoming clients connect to this port
- Once server **accepts** a connection, it talks to client on a **different “ephemeral” port**



# Echo Demo

- Try to connect two clients to the same server.
- What happens?
  - Second client has to wait for first client to finish!
  - Server doesn't even accept second client's connection
  - Where/why are we getting stuck?
- Because we're stuck in echo() talking to the first client, echoserver can't handle any more clients
- Solution: multi-threading

# Echo Server Multithreaded

- How might we make this server multithreaded?  
(Don't look at echoserver\_t.c)

```
while (1) {
 // Allocate space on the stack for client info
 client_info client_data;
 client_info *client = &client_data;

 // Initialize the length of the address
 client->addr = sizeof(client->addr);

 // Accept() will block until a client connects to the port
 client->connfd = Accept(listenfd,
 (SA *) &client->addr, &client->addr);
}

// Connection is established; echo to client
echo(client);
```

# Echo Server Multithreaded

- **echoserver\_t.c isn't too different from echoserver.c**
  - To see the changes: `diff echoserver.c echoserver\_t.c`
- **Making your proxy multithreaded will be very similar**
- **However, don't underestimate the difficulty of addressing race conditions between threads!**
  - Definitely the hardest part of proxylab
  - More on this next time...